

```
import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.tree import DecisionTreeClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

```
from sklearn.cluster import KMeans, AgglomerativeClustering, DBSCANs

from mlxtend.frequent_patterns import apriori, association_rules

from google.colab import files
uploaded = files.upload()
```

No file chosen

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

```
df = pd.read_csv("placement_data.csv")
df.head()
```

	student_id	gender	degree	stream	internship_score	cgpa	placed
0	1	Male	B.Tech	IT	80	8.5	Yes
1	2	Female	B.Tech	CS	70	8.0	Yes
2	3	Male	B.E	EXTC	50	6.5	No
3	4	Female	B.E	IT	90	9.1	Yes
4	5	Male	B.Tech	MECH	40	6.0	No

```
# Check missing values
print(df.isnull().sum())
```

```
student_id      0
gender          0
degree          0
stream          0
internship_score 0
cgpa            0
placed          0
dtype: int64
```

```
df.fillna(df.mean(numeric_only=True), inplace=True)
```

```
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
for col in df.select_dtypes(include=['object']).columns:
    df[col] = le.fit_transform(df[col])
```

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
scaled_data = scaler.fit_transform(df)
df_scaled = pd.DataFrame(scaled_data, columns=df.columns)
df_scaled.head()
```

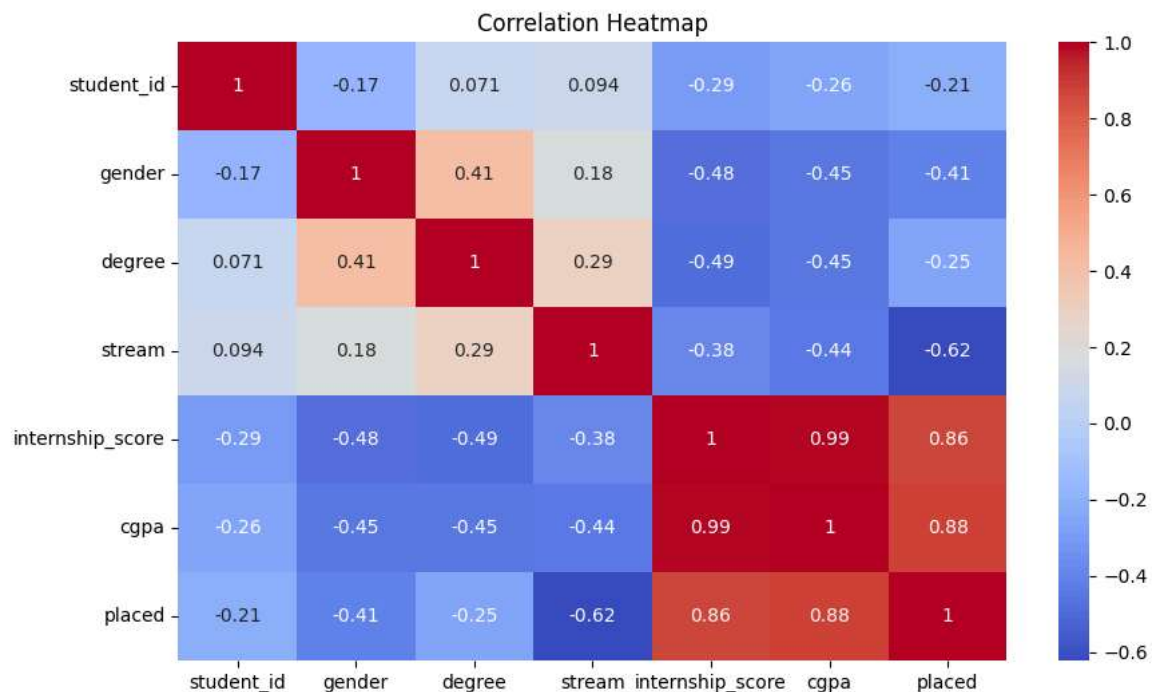
	student_id	gender	degree	stream	internship_score	cgpa	placed
0	-1.566699	1.0	0.816497	0.538816	0.904534	0.831670	0.816497
1	-1.218544	-1.0	0.816497	-1.257237	0.301511	0.359130	0.816497
2	-0.870388	1.0	-1.224745	-0.359211	-0.904534	-1.058490	-1.224745
3	-0.522233	-1.0	-1.224745	0.538816	1.507557	1.398718	0.816497
4	-0.174078	1.0	0.816497	1.436842	-1.507557	-1.531030	-1.224745

```
print(df.describe()) # summary (mean, min, max, etc.)
```

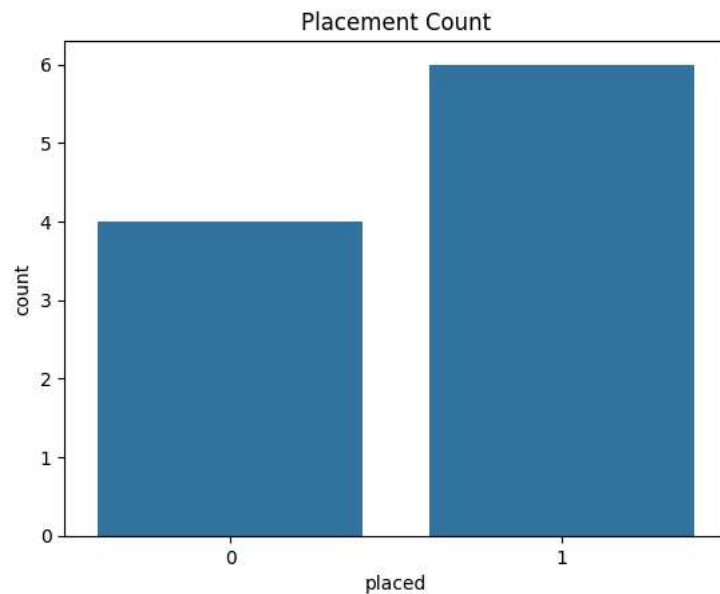
	student_id	gender	degree	stream	internship_score
count	10.00000	10.00000	10.00000	10.00000	10.00000
mean	5.50000	0.50000	0.60000	1.40000	65.00000
std	3.02765	0.52704	0.51639	1.17378	17.48014
min	1.00000	0.00000	0.00000	0.00000	40.00000
25%	3.25000	0.00000	0.00000	0.25000	51.25000
50%	5.50000	0.50000	1.00000	1.50000	65.00000
75%	7.75000	1.00000	1.00000	2.00000	78.75000
max	10.00000	1.00000	1.00000	3.00000	90.00000

	cgpa	placed
count	10.00000	10.00000
mean	7.62000	0.60000
std	1.11534	0.51639
min	6.00000	0.00000
25%	6.67500	0.00000
50%	7.75000	1.00000
75%	8.47500	1.00000
max	9.10000	1.00000

```
plt.figure(figsize=(10,6))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm')
plt.title("Correlation Heatmap")
plt.show()
```



```
sns.countplot(x='placed', data=df)
plt.title("Placement Count")
plt.show()
```




```
X = df.drop('placed', axis=1)
y = df['placed']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
dt = DecisionTreeClassifier()
dt.fit(X_train, y_train)
y_pred_dt = dt.predict(X_test)
```

```
cm = confusion_matrix(y_test, y_pred_dt)
plt.figure(figsize=(6,4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix Heatmap')
plt.show()
```



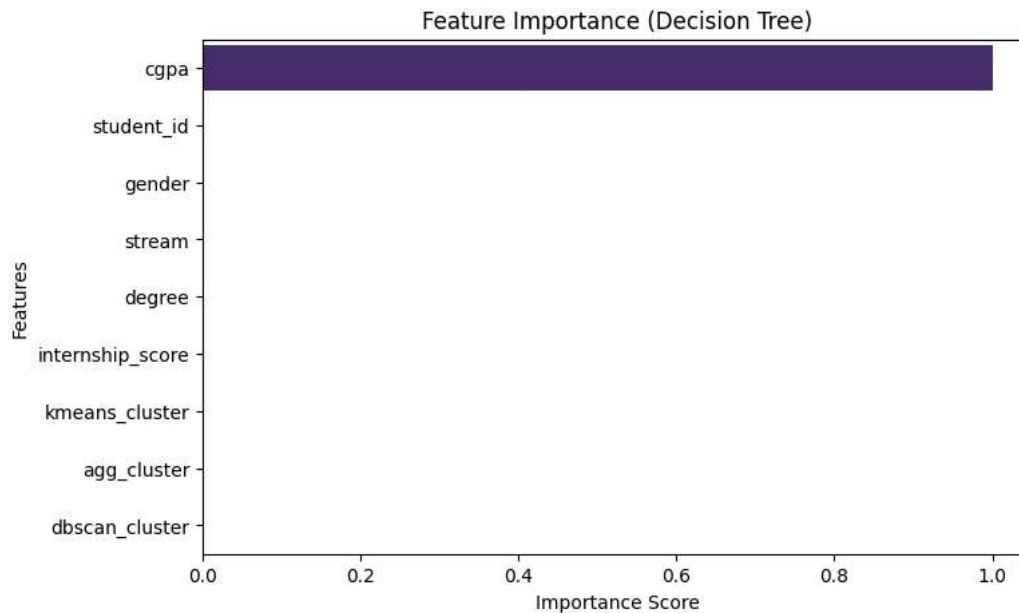
```
importances = dt.feature_importances_
feature_names = X.columns
feat_imp = pd.Series(importances, index=feature_names).sort_values(ascending=False)

plt.figure(figsize=(8,5))
sns.barplot(x=feat_imp.values, y=feat_imp.index, palette='viridis')
plt.title('Feature Importance (Decision Tree)')
plt.xlabel('Importance Score')
plt.ylabel('Features')
plt.show()
```

```
/tmp/ipython-input-916379644.py:6: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend`

```
sns.barplot(x=feat_imp.values, y=feat_imp.index, palette='viridis')
```



```
nb = GaussianNB()
nb.fit(X_train, y_train)
y_pred_nb = nb.predict(X_test)

print("Naive Bayes Accuracy:", accuracy_score(y_test, y_pred_nb))
print(confusion_matrix(y_test, y_pred_nb))
print(classification_report(y_test, y_pred_nb))
```

Naive Bayes Accuracy: 1.0

```
[[2]]
```

	precision	recall	f1-score	support
1	1.00	1.00	1.00	2
accuracy			1.00	2
macro avg	1.00	1.00	1.00	2
weighted avg	1.00	1.00	1.00	2

```
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:407: UserWarning: A single label was found in 'y_true' and
warnings.warn(
```

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.cluster import KMeans, AgglomerativeClustering, DBSCAN
```

```
from google.colab import files
uploaded = files.upload()
```

No file chosen

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

```
df = pd.read_csv("placement_data.csv")
print("Dataset loaded successfully!")
df.head()
```

Dataset loaded successfully!

	student_id	gender	degree	stream	internship_score	cgpa	placed
0	1	Male	B.Tech	IT	80	8.5	Yes
1	2	Female	B.Tech	CS	70	8.0	Yes
2	3	Male	B.E	EXTC	50	6.5	No
3	4	Female	B.E	IT	90	9.1	Yes
4	5	Male	B.Tech	MECH	40	6.0	No

```
display(df.head())
```

	student_id	gender	degree	stream	internship_score	cgpa	placed
0	1	Male	B.Tech	IT	80	8.5	Yes
1	2	Female	B.Tech	CS	70	8.0	Yes
2	3	Male	B.E	EXTC	50	6.5	No
3	4	Female	B.E	IT	90	9.1	Yes
4	5	Male	B.Tech	MECH	40	6.0	No

```
df.fillna(df.mean(numeric_only=True), inplace=True)
```

```
display(df.head())
```

```
le = LabelEncoder()
for col in df.select_dtypes(include=['object']).columns:
    df[col] = le.fit_transform(df[col])
```


```
scaler = StandardScaler()
scaled_data = scaler.fit_transform(df)
df_scaled = pd.DataFrame(scaled_data, columns=df.columns)
```

```
print(df.columns)
```

```
Index(['student_id', 'gender', 'degree', 'stream', 'internship_score', 'cgpa',
      'placed', 'kmeans_cluster', 'agg_cluster', 'dbscan_cluster'],
      dtype='object')
```

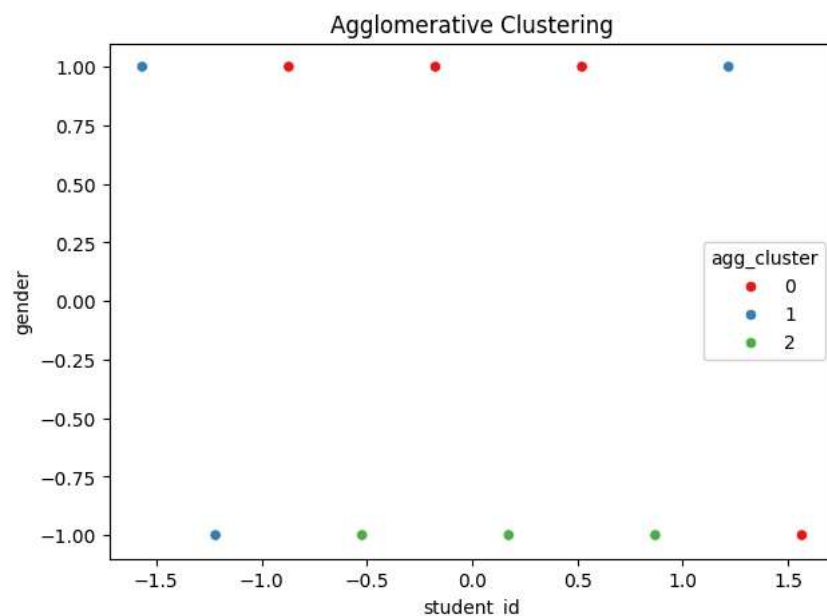
```
kmeans = KMeans(n_clusters=3, random_state=42)
df['kmeans_cluster'] = kmeans.fit_predict(df_scaled)
```

```
plt.figure(figsize=(7,5))
sns.scatterplot(x=df_scaled.iloc[:,0], y=df_scaled.iloc[:,1], hue=df['kmeans_cluster'], palette='Set2')
plt.title("K-Means Clustering")
plt.show()
```



```
agg = AgglomerativeClustering(n_clusters=3)
df['agg_cluster'] = agg.fit_predict(df_scaled)
```

```
plt.figure(figsize=(7,5))
sns.scatterplot(x=df_scaled.iloc[:,0], y=df_scaled.iloc[:,1], hue=df['agg_cluster'], palette='Set1')
plt.title("Agglomerative Clustering")
plt.show()
```



```
dbscan = DBSCAN(eps=1.5, min_samples=3)
df['dbscan_cluster'] = dbscan.fit_predict(df_scaled)
```

```
plt.figure(figsize=(7,5))
sns.scatterplot(x=df_scaled.iloc[:,0], y=df_scaled.iloc[:,1], hue=df['dbscan_cluster'], palette='Set3')
plt.title("DBSCAN Clustering")
plt.show()
```



```
print(df[['kmeans_cluster', 'agg_cluster', 'dbscan_cluster']].head())
```

	-1.5	-1.0	-0.5	0.0	0.5	1.0	1.5
	kmeans_cluster	agg_cluster	dbscan_cluster	student_id			
0	2	1	-1				
1	2	1	-1				
2	1	0	-1				
3	0	2	-1				
4	1	0	-1				

```
# Use existing column names
df_apriori = df[['gender', 'degree', 'placed']] # replace degree_t with actual column name
df_apriori = pd.get_dummies(df_apriori)

# Find frequent itemsets
frequent_itemsets = apriori(df_apriori, min_support=0.2, use_colnames=True)

# Generate rules
rules = association_rules(frequent_itemsets, metric='lift', min_threshold=1)
print(rules[['antecedents', 'consequents', 'support', 'confidence', 'lift']])
```

	antecedents	consequents	support	confidence	lift
0	(gender)	(degree)	0.4	0.800000	1.333333
1	(degree)	(gender)	0.4	0.666667	1.333333
2	(gender, placed)	(degree)	0.2	1.000000	1.666667
3	(degree, placed)	(gender)	0.2	0.666667	1.333333
4	(gender)	(degree, placed)	0.2	0.400000	1.333333
5	(degree)	(gender, placed)	0.2	0.333333	1.666667

```
df_apriori = df[['gender', 'degree', 'placed']]
df_apriori = pd.get_dummies(df_apriori)
```

```
frequent_itemsets = apriori(df_apriori, min_support=0.2, use_colnames=True)
rules = association_rules(frequent_itemsets, metric='lift', min_threshold=1)
print(rules[['antecedents', 'consequents', 'support', 'confidence', 'lift']])
```